

BDA2 - Spark SQL

Roshni Sundaramurthy (rossu809) and Yusur Al-Mter (yusal621)

May 12, 2019

Contents

1	Question 1.	2
1.1	Query	2
1.2	Extract	3
2	Question 2.	4
2.1	Query using APIs	4
2.2	Regular query	5
2.3	Extract	5
3	Question 3.	6
3.1	Query	6
3.2	Extract	6
4	Question 4.	7
4.1	Query	7
4.2	Extract	7
5	Question 5.	8
5.1	Query	8
5.2	Extract	8
6	Question 6.	9
6.1	Query	9
6.2	Extract	10

Exercises BDA 2 using Spark SQL

Imports

```
from pyspark import SparkContext
from pyspark.sql import SQLContext, Row
from pyspark.sql import functions as F
```

```
##### Temperature file #####

# Load a text file and convert each line to a tuple.
sc = SparkContext()
rdd = sc.textFile("/user/x_rossu/bdlab1/temperature-readings.csv")

sqlContext = SQLContext(sc)
parts = rdd.map(lambda l: l.split(","))
tempReadingsRow = parts.map(lambda p: (p[0], p[1], int(p[1].split("-")[0]),
int(p[1].split("-")[1]), p[2], float(p[3]), p[4] ))

# Specifying the schema programatically and registering the DataFrame as a table
tempReadingsString = ["station", "date", "year", "month", "time", "value",
"quality"]
```

```

# Apply the schema to the RDD.
schemaTempReadings = sqlContext.createDataFrame(tempReadingsRow,
tempReadingsString)

# Register the DataFrame as a table.
schemaTempReadings.registerTempTable("tempReadingsTable")

##### Precipitation file #####

# Data importing
rdd = sc.textFile("/user/x_rossu/bdlab1/precipitation-readings.csv")
parts = rdd.map(lambda l: l.split(";"))
precipReadingsRow = parts.map(lambda p: (int(p[0]), p[1], int(p[1].split("-")[0]),
int(p[1].split("-")[1]), p[2], float(p[3]), p[4]))

precipReadingsString = ["station", "date", "year", "month", "time", "precipitation", "quality"]

# Apply the schema to the RDD
schemaPrecipReadings = sqlContext.createDataFrame(precipReadingsRow, precipReadingsString)

# Register the DataFrame as a table
schemaPrecipReadings.registerTempTable("precipReadingsTable")

##### Station Ostergotland file #####

rdd = sc.textFile("C:/Users/roshn/Desktop/Bigdata/stations-Ostergotland.csv")
parts = rdd.map(lambda l: l.split(";"))
StationReadingsRow = parts.map(lambda p: (int(p[0]), p[1] ))

StationReadingsString = ["station", "stn_name"]

# Apply the schema to the RDD
schemaStationReadings = sqlContext.createDataFrame(StationReadingsRow, StationReadingsString)

# Register the DataFrame as a table
schemaStationReadings.registerTempTable("StationReadingsTable")

# Can run queries now

```

1 Question 1.

year, station with the max, maxValue ORDER BY maxValuE DESC
year, station with the min, minValuE ORDER BY minValuE DESC

1.1 Query

```
# Running SQL queries - API methods
```

```

#max temp

schemaTempReadingsMax = schemaTempReadings.filter(schemaTempReadings['year'] >= 1950)
                        .filter(schemaTempReadings.year <= 2014)
schemaTempReadingsMax = schemaTempReadingsMax.groupBy('year').agg(F.max('value').alias('value'))
                        .join(schemaTempReadings, ['year', 'value']).drop_duplicates(['year'])
                        .select(['year', 'station', 'value']).orderBy(['value'],ascending=[0])
schemaTempReadingsMax.show()

# min temp

schemaTempReadingsMin = schemaTempReadings.filter(schemaTempReadings['year'] >= 1950)
                        .filter(schemaTempReadings.year <= 2014)
schemaTempReadingsMin = schemaTempReadingsMin.groupBy('year').agg(F.min('value').alias('value'))
                        .join(schemaTempReadings, ['year', 'value']).dropDuplicates(['year'])
                        .select(['year', 'station', 'value']).orderBy(['value'],ascending=[0])
schemaTempReadingsMin.show()

```

1.2 Extract

Lowest and highest temperatures measured each year for the period 1950-2014:

```

# Max-temperatures

+-----+-----+-----+
|year|station|value|
+-----+-----+-----+
|1975| 86200| 36.1|
|1992| 63600| 35.4|
|1994| 117160| 34.7|
|2010| 75250| 34.4|
|2014| 96560| 34.4|
|1989| 63050| 33.9|
|1982| 94050| 33.8|
|1968| 137100| 33.7|
|1966| 151640| 33.5|
|1983| 98210| 33.3|
|2002| 78290| 33.3|
|1970| 103080| 33.2|
|1986| 76470| 33.2|
|2000| 62400| 33.0|
|1956| 145340| 33.0|
|1959| 65160| 32.8|
|2006| 75240| 32.7|
|1991| 137040| 32.7|
|1988| 102540| 32.6|
|2011| 172770| 32.5|
+-----+-----+-----+
only showing top 20 rows

# Min-temperaures

```

```

+-----+-----+-----+
|year|station|value|
+-----+-----+-----+
|1990| 166870|-35.0|
|1952| 192830|-35.5|
|1974| 179950|-35.6|
|1954| 113410|-36.0|
|1992| 179960|-36.1|
|1975| 157860|-37.0|
|1972| 167860|-37.5|
|1995| 182910|-37.6|
|2000| 169860|-37.6|
|1957| 159970|-37.8|
|1989| 166870|-38.2|
|1983| 191900|-38.2|
|1953| 183760|-38.4|
|2009| 179960|-38.5|
|1993| 191900|-39.0|
|1984| 191900|-39.2|
|1991| 179960|-39.3|
|1973| 166870|-39.3|
|2008| 179960|-39.3|
|2005| 155790|-39.4|
+-----+-----+-----+
only showing top 20 rows

```

2 Question 2.

year, month, value ORDER BY value DESC
year, month, value ORDER BY value DESC

2.1 Query using APIs

```

schemaTempReadingscount = schemaTempReadings.filter(schemaTempReadings['value'] > 10)
                        .filter(schemaTempReadings['year'] >= 1950)
                        .filter(schemaTempReadings.year <= 2014)
schemaTempReadingscount = schemaTempReadingscount.groupBy('year', 'month').agg(F.count('value')
                        .alias('count')).orderBy(['count'], ascending=[0])
schemaTempReadingscount.show()

schemaTempReadingscount = schemaTempReadings.filter(schemaTempReadings['value'] > 10)
                        .filter(schemaTempReadings['year'] >= 1950)
                        .filter(schemaTempReadings.year <= 2014)
schemaTempReadingscount = schemaTempReadingscount.groupBy('year', 'month')
                        .agg(F.countDistinct('station')
                        .alias('count')).orderBy(['count'], ascending=[0])
schemaTempReadingscount.show()

```

2.2 Regular query

```
count = sqlContext.sql("SELECT year,month, count(value) as value FROM tempReadingsTable
                        WHERE year>=1950 and year<=2014 and value>=10.0
                        GROUP BY year, month ORDER BY value DESC")

count.show()

count = sqlContext.sql("SELECT year,month, count(tmp) as value
                        FROM (SELECT year, month, station, count(value) as tmp
                              FROM tempReadingsTable
                              WHERE year>=1950 and year<=2014 and value>=10.0
                              GROUP BY year, month, station)
                        GROUP BY year, month ORDER BY value DESC")

count.show()
```

2.3 Extract

Number of readings above 10 degrees for each month in the period of 1950-2014:

year	month	value
2014	7	147910
2011	7	147060
2010	7	143860
2012	7	138166
2013	7	134297
2009	7	133570
2011	8	133483
2009	8	129007
2013	8	128920
2003	7	128360
2002	7	128354
2006	8	128039
2008	7	127627
2002	8	126495
2011	6	126084
2012	8	125947
2005	7	125651
2006	7	125192
2010	8	125135
2014	8	125006

Distinct readings above 10 degrees for each month in the period of 1950-2014:

year	month	count
1972	10	378
1973	5	377
1973	6	377
1973	9	376
1972	8	376
1972	6	375
1972	9	375

```
|1972|    5|  375|
|1971|    8|  375|
|1972|    7|  374|
|1971|    9|  374|
|1971|    6|  374|
|1971|    5|  373|
|1973|    8|  373|
|1974|    8|  372|
|1974|    6|  372|
|1973|    7|  370|
|1970|    8|  370|
|1974|    9|  370|
|1971|    7|  370|
+-----+
```

3 Question 3.

year, month, station, avgMonthlyTemperature ORDER BY avgMonthlyTemperature DESC

3.1 Query

```
schemayearReadings = schemaTempReadings.filter(schemaTempReadings.year >= 1960)
                                .filter(schemaTempReadings.year <= 2014)
avg_temp = schemayearReadings.groupBy(['station','date', 'year', 'month'])
                                .agg(F.max('value').alias('max_temp')
                                    ,F.min('value').alias('min_temp'))
avg_temp = avg_temp.select(['station', 'year', 'month', 'min_temp', 'max_temp'])
avg_temp = avg_temp.withColumn('sum', avg_temp['min_temp'] + avg_temp['max_temp'])
                                .groupBy(['station', 'year', 'month']).agg(F.avg('sum').alias('value'))
avg_temp = avg_temp.withColumn('avg_temp', avg_temp['value']/2)
                                .select(['year', 'month', 'station', 'avg_temp'])
                                .orderBy('avg_temp', ascending = [0])

avg_temp.show()
```

3.2 Extract

Average monthly temperature for each available station in Sweden in the period of 1960-2014:

```
+-----+
|year|month|station|          avg_temp|
+-----+
|2014|    7| 96000|          26.3|
|1994|    7| 96550|23.071052631578947|
|1983|    8| 54550|          23.0|
|1994|    7| 78140|22.970967741935485|
|1994|    7| 85280| 22.87258064516129|
|1994|    7| 75120|22.858064516129037|
|1994|    7| 65450| 22.85645161290323|
|1994|    7| 96000|22.808064516129033|
```

1994	7	95160	22.764516129032256
1994	7	86200	22.71129032258065
2002	8	78140	22.7
1994	7	76000	22.698387096774194
1997	8	78140	22.666129032258066
1994	7	105260	22.65967741935484
1975	8	54550	22.642857142857142
2006	7	76530	22.598387096774193
1994	7	86330	22.54838709677419
2006	7	75120	22.52741935483871
1994	7	54300	22.469354838709677
2006	7	78140	22.458064516129028

+-----+

4 Question 4.

station, maxTemp, maxDailyPrecipitation ORDER BY station DESC

4.1 Query

```
# Temp
max_temp = schemaTempReadings.groupBy('station').agg(F.max('value').alias('maxtempvalue'))
max_temp = max_temp.filter(max_temp['maxtempvalue'] >= 25)
                    .filter(max_temp['maxtempvalue'] <= 30)
                    .orderBy(['station'], ascending=[0])
max_temp.show()

# Precipitation
max_precip = schemaPrecipReadings.groupBy('station', 'date')
                    .agg(F.sum('precipitation').alias('total_precip'))
max_precip = max_precip.select('station', 'total_precip').groupBy('station')
                    .agg(F.max('total_precip').alias('maxprecipvalue'))
max_precip = max_precip.filter(max_precip['maxprecipvalue'] >= 100)
                    .filter(max_precip['maxprecipvalue'] <= 200)
                    .orderBy(['station'], ascending=[0])
max_precip.show()
```

4.2 Extract

Stations with their associated maximum measured temperatures and maximum measured daily precipitation:

```
# Max_Temperatures
|station|maxtempvalue|
+-----+
| 99450|      26.0|
| 99280|      25.5|
| 99270|      27.8|
| 99090|      27.6|
| 98610|      29.2|
| 98170|      27.0|
```

98140	26.4
96600	26.2
96370	30.0
96220	30.0
95640	29.5
95620	29.8
95380	26.1
95230	29.9
94660	29.3
94450	29.5
94190	30.0
93640	30.0
93250	28.6
91620	30.0

Max_precipitation

station	maxprecipvalue
97510	103.99999999999999
75250	101.8
71420	106.3
52350	101.6

5 Question 5.

year, month, avgMonthlyPrecipitation ORDER BY year DESC, month DESC

5.1 Query

```
precip = schemaStationReadings.join(schemaPrecipReadings, ['station'])
precip = precip.select('station', 'year', 'month', 'precipitation')
               .filter(precip['year'] >= 1993).filter(precip.year <= 2016)
precip = precip.groupBy('station', 'year', 'month')
               .agg(F.sum('precipitation').alias('precip'))
precip = precip.select('year', 'month', 'precip')
precip = precip.groupBy('year', 'month').agg(F.avg('precip').alias('avg_precip'))
               .orderBy(['year', 'month'], ascending = False)

precip.show()
```

5.2 Extract

Average monthly precipitation for the Östergötland region for the period 1993-2016:

year	month	avg_precip
2016	7	0.0
2016	6	47.662499999999994
2016	5	29.250000000000004
2016	4	26.900000000000001
2016	3	19.962500000000002
2016	2	21.5625
2016	1	22.325
2015	12	28.925000000000004
2015	11	63.887500000000001
2015	10	2.2625
2015	9	101.29999999999997
2015	8	26.9875
2015	7	119.09999999999997
2015	6	78.662500000000001
2015	5	93.22499999999998
2015	4	15.337499999999999
2015	3	42.612500000000001
2015	2	24.824999999999996
2015	1	59.112500000000026
2014	12	35.462500000000001

6 Question 6.

year, month, difference ORDER BY year DESC, month DESC

6.1 Query

```
temp = schemaStationReadings.join(schemaTempReadings, ['station'])
temp = temp.select('station', 'date', 'year', 'month', 'value')
        .filter(temp['year'] >= 1950).filter(temp['year'] <= 2014)
temp = temp.groupBy(['station', 'date', 'year', 'month'])
        .agg(F.min('value').alias('min'), F.max('value').alias('max'))
temp = temp.select(['station', 'year', 'month', 'min', 'max'])
        .withColumn('sum_temp', temp['min'] + temp['max'])
temp = temp.groupBy(['station', 'year', 'month']).agg(F.avg('sum_temp').alias('value'))
temp = temp.withColumn('avg_temp', temp.value/2)
avg_temp = temp.select(['year', 'month', 'station', 'avg_temp'])
        .groupBy(['year', 'month']).agg(F.avg('avg_temp').alias('avg_temp'))
long_temp = avg_temp.filter(avg_temp['year'] >= 1950).filter(avg_temp['year'] <= 1980)
long_temp = long_temp.groupBy('month').agg(F.avg('avg_temp')
        .alias('long_temp')).join(avg_temp, 'month')
long_temp = long_temp.withColumn('Compared_Temperature'
        , long_temp['avg_temp'] - long_temp['long_temp'])
long_temp = long_temp.select(['year', 'month', 'Compared_Temperature'])
long_temp = long_temp.orderBy(['year', 'month'], ascending = False)
long_temp.show()
```

```
# if plot
#long_temp.rdd.coalesce(1).saveAsTextFile("comp_temp")
```

6.2 Extract

Compared average monthly temperature in the period 1950-2014 with long-term monthly averages in the period of 1950-1980:

```
+-----+
|year|month|Compared_temperature|
+-----+
|2014| 12| 0.8238893537957012|
|2014| 11| 2.0635396726928987|
|2014| 10| 1.5225549840378134|
|2014|  9| 0.06105818643721861|
|2014|  8| -0.6426470719706909|
|2014|  7| 2.0939107758930824|
|2014|  6| -1.8073686197315197|
|2014|  5| 0.26719065014069976|
|2014|  4| 2.0661931589915454|
|2014|  3| 4.486748343574566|
|2014|  2| 5.420311314566043|
|2014|  1| 0.9325880207201753|
|2013| 12| 3.8796729966761174|
|2013| 11| 0.9342517939050206|
|2013| 10| 0.7529068901961731|
|2013|  9| -1.00505751604212|
|2013|  8| -0.31464120686804975|
|2013|  7| 0.008280277359357768|
|2013|  6| -0.5441868015497029|
|2013|  5| 1.573920855419292|
+-----+
```